

```
In [245]: import pandas as pd
```

```
In [246]: data=pd.read_csv("/home/placement/Desktop/csv/fiat500.csv")
```

```
In [247]: data.describe()
```

Out[247]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>count</b>	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
<b>mean</b>	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
<b>std</b>	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
<b>min</b>	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
<b>25%</b>	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
<b>50%</b>	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
<b>75%</b>	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
<b>max</b>	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

```
In [248]: data.head()
```

Out[248]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>0</b>	1	lounge	51	882	25000	1	44.907242	8.611560	8900
<b>1</b>	2	pop	51	1186	32500	1	45.666359	12.241890	8800
<b>2</b>	3	sport	74	4658	142228	1	45.503300	11.417840	4200
<b>3</b>	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
<b>4</b>	5	pop	73	3074	106880	1	41.903221	12.495650	5700

In [249]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    1538 non-null  int64
1   model                 1538 non-null  object
2   engine_power          1538 non-null  int64
3   age_in_days           1538 non-null  int64
4   km                    1538 non-null  int64
5   previous_owners       1538 non-null  int64
6   lat                   1538 non-null  float64
7   lon                   1538 non-null  float64
8   price                 1538 non-null  int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

```
In [250]: data1=data.drop(['ID','lat','lon'],axis=1)  
data1
```

Out[250]:

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [251]: data1.shape
```

Out[251]: (1538, 6)

```
In [252]: data1=pd.get_dummies(data1)
data1
```

Out[252]:

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...	...	...	...	...	...	...	...	...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [253]: data1.shape
```

Out[253]: (1538, 8)

```
In [254]: y=data1['price']
x=data1.drop(['price'],axis=1)
```

In [255]:

y

Out[255]:

0	8900
1	8800
2	4200
3	6000
4	5700

...

1533	5200
1534	4600
1535	7500
1536	5990
1537	7900

Name: price, Length: 1538, dtype: int64

In [256]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random_state=42)
```

In [257]:

x\_test.shape

Out[257]:

(508, 7)

In [258]:

x\_train.head(5)

Out[258]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0

```
In [259]: y_test.head(5)
```

```
Out[259]: 481      7900  
          76      7900  
          1502    9400  
          669    8500  
          1409    9700  
          Name: price, dtype: int64
```

```
In [260]: y_train.head(5)
```

```
Out[260]: 527      9990  
          129      9500  
          602      7590  
          331      8750  
          323      9100  
          Name: price, dtype: int64
```

```
In [261]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
```

```
alpha = [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20, 30]
```

```
ridge = Ridge()
```

```
parameters = {'alpha': alpha}
```

```
ridge_regressor = GridSearchCV(ridge, parameters)
```

```
ridge_regressor.fit(x_train, y_train)
```


```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=5.56109e-26): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.70876e-26): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=6.91585e-23): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.08003e-23): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.01022e-23): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.57959e-23): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.24161e-23): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=6.92759e-21): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.09091e-21): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
```

```

Ill-conditioned matrix (rcond=7.02112e-21): result may not be accurate.
return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.57414e-21): result may not be accurate.
return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.23284e-21): result may not be accurate.
return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=6.9277e-17): result may not be accurate.
return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.09099e-17): result may not be accurate.
return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.02123e-17): result may not be accurate.
return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.57407e-17): result may not be accurate.
return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.23274e-17): result may not be accurate.
return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T

```

Out[261]:



In [262]: `ridge_regressor.best_params_`

Out[262]: `{'alpha': 30}`

```

In [263]: ridge=Ridge(alpha=30)
          ridge.fit(x_train,y_train)
          y_pred_ridge=ridge.predict(x_train)

```



```
In [264]: from sklearn.metrics import r2_score  
r2_score(y_test,ypred)
```

Out[264]: 0.8415526986865394

```
In [265]: from sklearn.metrics import mean_squared_error  
Ridge_error=mean_squared_error(ypred,y_test)  
Ridge_error
```

Out[265]: 581887.727391353

```
In [266]: results=pd.DataFrame(columns=['Actual','predicted'])
results['Actual']=y_test
results['predicted']=ypred
results=results.reset_index()
results['ID']=results.index
results.head(15)
```

Out[266]:

	index	Actual	predicted	ID
0	481	7900	5867.650338	0
1	76	7900	7133.701423	1
2	1502	9400	9866.357762	2
3	669	8500	9723.288745	3
4	1409	9700	10039.591012	4
5	1414	9900	9654.075826	5
6	1089	9900	9673.145630	6
7	1507	9950	10118.707281	7
8	970	10700	9903.859527	8
9	1198	8999	9351.558284	9
10	1088	9890	10434.349636	10
11	576	7990	7732.262557	11
12	965	7380	7698.672401	12
13	1488	6800	6565.952404	13
14	1432	8900	9662.901035	14

```
In [272]: data=data.loc[(data.model=='lounge')]  
data
```


Out[272]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
6	7	lounge	51	731	11600	1	44.907242	8.611560	10750
7	8	lounge	51	1521	49076	1	41.903221	12.495650	9190
11	12	lounge	51	366	17500	1	45.069679	7.704920	10990
...	...	...	...	...	...	...	...	...	...
1528	1529	lounge	51	2861	126000	1	43.841980	10.515310	5500
1529	1530	lounge	51	731	22551	1	38.122070	13.361120	9900
1530	1531	lounge	51	670	29000	1	45.764648	8.994500	10800
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990

1094 rows × 9 columns

```
In [270]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='ID', y='Actual', data=results.head(50))
sns.lineplot(x='ID', y='predicted', data=results.head(50))
```

Out[270]: <Axes: xlabel='ID', ylabel='Actual'>



In [ ]:

In [ ]: